



Suki Lau

[Follow](#)

Enthusiastic learner in machine learning, deep learning, natural language processing.

Jul 29, 2017 · 6 min read

## Learning Rate Schedules and Adaptive Learning Rate Methods for Deep Learning

When training deep neural networks, it is often useful to reduce learning rate as the training progresses. This can be done by using pre-defined **learning rate schedules** or **adaptive learning rate methods**. In this article, I train a convolutional neural network on [CIFAR-10](#) using differing learning rate schedules and adaptive learning rate methods to compare their model performances.

### Learning Rate Schedules

Learning rate schedules seek to adjust the learning rate during training by reducing the learning rate according to a pre-defined schedule. Common learning rate schedules include **time-based decay**, **step decay** and **exponential decay**. For illustrative purpose, I construct a convolutional neural network trained on [CIFAR-10](#), using stochastic gradient descent (SGD) optimization algorithm with different learning rate schedules to compare the performances.

#### Constant Learning Rate

Constant learning rate is the default learning rate schedule in [SGD optimizer](#) in Keras. Momentum and decay rate are both set to zero by default. It is tricky to choose the right learning rate. By experimenting with range of learning rates in our example, `lr=0.1` shows a relative good performance to start with. This can serve as a baseline for us to experiment with different learning rate strategies.

```
keras.optimizers.SGD(lr=0.1, momentum=0.0, decay=0.0,  
nesterov=False)
```




Fig 1 : Constant Learning Rate

## Time-Based Decay

The mathematical form of time-based decay is  $lr = lr_0 / (1 + kt)$  where  $lr$ ,  $k$  are hyperparameters and  $t$  is the iteration number. Looking into the [source code](#) of Keras, the SGD optimizer takes `decay` and `lr` arguments and update the learning rate by a decreasing factor in each epoch.

```
lr *= (1. / (1. + self.decay * self.iterations))
```

Momentum is another argument in SGD optimizer which we could tweak to obtain faster convergence. Unlike classical SGD, momentum method helps the parameter vector to build up velocity in any direction with constant gradient descent so as to prevent oscillations. A typical choice of momentum is between 0.5 to 0.9.

SGD optimizer also has an argument called `nesterov` which is set to false by default. Nesterov momentum is a different version of the momentum method which has stronger theoretical converge guarantees for convex functions. In practice, it works slightly better than standard momentum.

In Keras, we can implement time-based decay by setting the initial learning rate, decay rate and momentum in the SGD optimizer.

```

learning_rate = 0.1
decay_rate = learning_rate / epochs
momentum = 0.8
sgd = SGD(lr=learning_rate, momentum=momentum,
          decay=decay_rate, nesterov=False)

```




Fig 2 : Time-based Decay Schedule

## Step Decay

Step decay schedule drops the learning rate by a factor every few epochs. The mathematical form of step decay is :

$$lr = lr_0 * drop^{\text{floor}(\text{epoch} / \text{epochs\_drop})}$$

A typical way is to drop the learning rate by half every 10 epochs. To implement this in Keras, we can define a step decay function and use `LearningRateScheduler` callback to take the step decay function as argument and return the updated learning rates for use in SGD optimizer.

```

def step_decay(epoch):
    initial_lrate = 0.1
    drop = 0.5
    epochs_drop = 10.0
    lrate = initial_lrate * math.pow(drop,
                                     math.floor((1+epoch)/epochs_drop))
    return lrate

```

```
lr_rate = LearningRateScheduler(step_decay)
```

As a digression, a callback is a set of functions to be applied at given stages of the training procedure. We can use callbacks to get a view on internal states and statistics of the model during training. In our example, we create a custom callback by extending the base class `keras.callbacks.Callback` to record loss history and learning rate during the training procedure.

```
class LossHistory(keras.callbacks.Callback):
    def on_train_begin(self, logs={}):
        self.losses = []
        self.lr = []

    def on_epoch_end(self, batch, logs={}):
        self.losses.append(logs.get('loss'))
        self.lr.append(step_decay(len(self.losses)))
```

Putting everything together, we can pass a callback list consisting of `LearningRateScheduler` callback and our custom callback to fit the model. We can then visualize the learning rate schedule and the loss history by accessing `loss_history.lr` and `loss_history.losses`.

```
loss_history = LossHistory()
lr_rate = LearningRateScheduler(step_decay)
callbacks_list = [loss_history, lr_rate]

history = model.fit(X_train, y_train,
                    validation_data=(X_test, y_test),
                    epochs=epochs,
                    batch_size=batch_size,
                    callbacks=callbacks_list,
                    verbose=2)
```




Fig 3a : Step Decay Schedule




Fig 3b : Step Decay Schedule

## Exponential Decay

Another common schedule is exponential decay. It has the mathematical form  $lr = lr_0 * e^{(-kt)}$ , where  $lr$ ,  $k$  are hyperparameters and  $t$  is the iteration number. Similarly, we can implement this by defining exponential decay function and pass it to `LearningRateScheduler`. In fact, any custom decay schedule can be implemented in Keras using this approach. The only difference is to define a different custom decay function.

```
def exp_decay(epoch):
    initial_lrate = 0.1
```

```
k = 0.1  
lrate = initial_lrate * exp(-k*t)  
return lrate
```

```
lrate = LearningRateScheduler(exp_decay)
```




Fig 4a : Exponential Decay Schedule




Fig 4b : Exponential Decay Schedule

Let us now compare the model accuracy using different learning rate schedules in our example.




Fig 5 : Comparing Performances of Different Learning Rate Schedules

## Adaptive Learning Rate Methods

The challenge of using learning rate schedules is that their hyperparameters have to be defined in advance and they depend heavily on the type of model and problem. Another problem is that the same learning rate is applied to all parameter updates. If we have sparse data, we may want to update the parameters in different extent instead.

Adaptive gradient descent algorithms such as **Adagrad**, **Adadelata**, **RMSprop**, **Adam**, provide an alternative to classical SGD. These per-parameter learning rate methods provide heuristic approach without requiring expensive work in tuning hyperparameters for the learning rate schedule manually.

In brief, **Adagrad** performs larger updates for more sparse parameters and smaller updates for less sparse parameters. It has good performance with sparse data and training large-scale neural networks. However, its monotonic learning rate usually proves too aggressive and stops learning too early when training deep neural networks. **Adadelata** is an extension of Adagrad that seeks to reduce its aggressive, monotonically decreasing learning rate. **RMSprop** adjusts the Adagrad method in a very simple way in an attempt to reduce its aggressive, monotonically decreasing learning rate. **Adam** is an update to the RMSprop optimizer which is like RMSprop with momentum.

In Keras, we can implement these adaptive learning algorithms easily using corresponding optimizers. It is usually recommended to leave the

hyperparameters of these optimizers at their default values (except `lr` sometimes).

```
keras.optimizers.Adagrad(lr=0.01, epsilon=1e-08, decay=0.0)
```

```
keras.optimizers.Adadelta(lr=1.0, rho=0.95, epsilon=1e-08,  
decay=0.0)
```

```
keras.optimizers.RMSprop(lr=0.001, rho=0.9, epsilon=1e-08,  
decay=0.0)
```

```
keras.optimizers.Adam(lr=0.001, beta_1=0.9, beta_2=0.999,  
epsilon=1e-08, decay=0.0)
```

Let us now look at the model performances using different adaptive learning rate methods. In our example, Adadelta gives the best model accuracy among other adaptive learning rate methods.




Fig 6 : Comparing Performances of Different Adaptive Learning Algorithms

Finally, we compare the performances of all the learning rate schedules and adaptive learning rate methods we have discussed.






Fig 7: Comparing Performances of Different Learning Rate Schedules and Adaptive Learning Algorithms

## Conclusion

In many examples I have worked on, adaptive learning rate methods demonstrate better performance than learning rate schedules, and they require much less effort in hyperparameter settings. We can also use `LearningRateScheduler` in Keras to create custom learning rate schedules which is specific to our data problem.

For further reading, Yoshua Bengio's paper provides very good practical recommendations for tuning learning rate for deep learning, such as how to set initial learning rate, mini-batch size, number of epochs and use of early stopping and momentum.

[Source code](#)

## References:

- [Practical Recommendations for Gradient-Based Training of Deep Architectures by Yoshua Bengio](#)
- [Convolutional Neural Networks for Visual Recognition](#)
- [Using Learning Rate Schedules for Deep Learning Models in Python with Keras](#)